

Robinhood PolicyEngine

Admin Guide

Temporary Filesystem Manager

v2.5.3

July 29th, 2014

<robinhood-support@lists.sourceforge.net>

Table of contents

1 Software presentation

4

1.1 Overview.....	4
1.2 Execution modes.....	5
1.3 Internal architecture overview.....	5

2 First steps with Robinhood

6

2.1 Compiling and installing.....	6
2.1.1 Install from RPM.....	6
2.1.2 Build and install from the source tarball.....	6
2.2 Robinhood service.....	7
2.3 Command line options.....	8
2.4 Signals.....	9
2.5 Creating the database.....	10
2.6 Enabling Lustre Changelogs.....	11

3 Writing configuration file

11

3.1 Syntax.....	11
3.2 Configuration template and default parameters.....	15
3.3 General parameters.....	16
3.4 Log and alerts parameters.....	16
3.5 File classes.....	18
3.6 Purge policies.....	18
3.7 Purge triggers.....	20
3.8 Purge parameters.....	22
3.9 Specifying 'rmdir' policy.....	23
3.10 'Rmdir' parameters.....	24
3.11 Periodic fileclass matching.....	24
3.12 Periodic information update when handling Lustre changelogs.....	24
3.13 Database parameters.....	25
3.14 Filesystem scan parameters.....	26
3.15 Lustre 2.x changelog parameters.....	27
[new 2.5] Lustre Changelog batching (optimization).....	28
[new 2.5] Reading Changelogs with Lustre DNE.....	28
[new 2.5] Dealing with old Changelogs bugs.....	29
3.16 Entry processor pipeline options.....	29

4 Reporting tool (rbh-report)

32

4.1 Overview.....	32
4.2 Command line.....	32
4.3 Reports.....	33

5 rbh-find

38

5.1 Overview.....	38
5.2 Syntax.....	38

5.3 Performance.....	39
6 rbh-du	
40	
6.1 Overview.....	40
6.2 Syntax.....	40
6.3 Performance.....	41
7 [new 2.5] rbh-diff	
41	
7.1 Overview.....	41
7.2 Syntax.....	41
7.3 Output overview.....	42
8 Web interface	
42	
8.1 Overview.....	42
8.2 Installation and configuration.....	43
9 Configuration and admin helper (rbh-config)	
43	
10 Performance tunings	
44	
10.1 Lustre tunings and workarounds.....	44
10.2 Linux kernel tunings.....	44
10.3 Optimize scanning vs reporting speed.....	45
10.4 Database tunings.....	45
10.5 DB operation batching vs. multi-threading.....	46
11 Getting help	
46	

1 Software presentation

1.1 Overview

“Robinhood Policy Engine” is an Open-Source software developed at CEA/DAM for monitoring and purging large temporary filesystems. It is designed in order to perform all its tasks in parallel, so it is particularly adapted for managing large file systems with millions of entries and petabytes of data. Moreover, it is Lustre capable i.e. it can monitor usage per OST and also purge files per OST and/or OST pools.

A specific mode of Robinhood also makes it possible to synchronize a cluster filesystem with a HSM by applying admin-defined migration and purge policies. However, this document only deals with temporary filesystem management purpose.

Temporary filesystem management mainly consists of the following aspects:

- Scan a large filesystem quickly, to build/update a list of candidate files in a database. This persistent storage ensures that a list of candidates is always available if a purge is needed (for freeing space in filesystem). Scanning is done using a parallel algorithm for best efficiency;
With Lustre v2, scanning is no longer needed thanks to Lustre MDT Changelog mechanism. An initial scan is still required to populate robinhood DB, then **Robinhood’s database is updated in soft real-time and you can use ‘rbh-find’ and ‘rbh-du’ commands to query your file system faster than ever!**
- Monitor filesystem and Lustre OST usage, in order to maintain them below a given threshold. If a storage unit exceeds the threshold, it then takes the most recent list of files and purges them in the order of their last access/modification time. It can also only purge the files of a given OST. Purge operations are also performed in parallel so it can quickly free disk space;
- Remove empty directories that have not be used for a long time, if the administrator wants so;
- Raise alerts, search for entries, generate accounting information and all kind of statistics about the filesystem...

All those actions are done according to very customizable policies. Thus, it makes it possible for you to preserve data of nice and responsible users, and penalize “abusers” and harmful behaviors... That’s why it is called Robinhood ;-)

Thanks to all of its features, Robinhood will help you to preserve the quality of service of your filesystem and avoid problematic situations.

1.2 Execution modes

Robinhood can be executed in 2 modes:

- As an ever-running daemon.
- As a “one-shot” command you can execute whenever you want.

In the daemon mode:

- Robinhood regularly refreshes its list of filesystem entries by scanning the filesystem it manages, and populating a database with this information.
If you run it on a Lustre v2 file system, it continuously reads MDT changelogs to update its database in soft real-time.
- It regularly monitors filesystem and OST usage, and purge entries whenever needed;
- It also regularly checks empty, unused directories, if you enabled this feature.

In one-shot mode, each action is made once, and then the program exits:

- It scans the filesystem once and update its database;
With Lustre v2, it reads MDT Changelogs currently stacked;
- It checks filesystem and OST usage, and purge entries only if needed;
- It checks empty directories, if you enabled this feature.

Note that you can use any combination of actions in daemon and one-shot mode. For example, you can make a daily scan of the filesystem (as a one-shot scan) and have a Robinhood daemon that only checks for filesystem and OST usage and purge entries when needed.

1.3 Internal architecture overview

Robinhood v2 uses a database engine for managing its list of filesystem entries, which offers a lot of benefits:

- It can manage larger filesystems, because the size of its list is not limited by the memory of the machine.
- The list it builds is persistent, so it can immediately purge filesystem entries, even after the daemon restarted. This also makes ‘one-shot’ runs possible.
- Scan and purge actions can be run on different nodes: no direct communication is needed between them; they only need to be database clients.
- Administrator can collect custom and complex statistics about filesystem content using a very standard language (SQL).

Filesystem entry processing is highly parallelized: a pool of threads is dedicated to scanning the namespace (readdir operations). Those threads then push listed entries to a pipeline. Then, other operations are performed asynchronously by another pool of threads:

- Getting attributes;
- Checking if entry is up-to-date in database;
- Getting stripe info if it is not already known;
- Checking alert rules and sending alerts;
- Insert/update the entry in the database.

Thanks to its Boolean expression engine, Robinhood policies and alert rules are flexible and easy to configure. A large range of attributes can be tested: name, path, type, owner, group, size, access or modification time, extended attributes, depth in namespace tree, number of entries in directory...

Various ways of triggering purges, to fit everyone's need:

- Purge triggers can be based on thresholds on OST usage or filesystem usage;
- Quota-like triggers can also be specified: if a given user or group exceeds a specified volume of data, then a purge is triggered on its files;
- Purge can be triggered when the inode count of a filesystem exceed a threshold.

2 First steps with Robinhood

2.1 *Compiling and installing*

2.1.1 Install from RPM

Pre-generated RPMs can be downloaded on sourceforge, for the following configurations:

- x86_64 architecture , RedHat 5/6 Linux family
- MySQL database 5.x
- Posix filesystems, Lustre 1.8, 2.1, 2.2, 2.3, 2.4, 2.5

Purpose specific RPM: `robinhood-tmpfs`

It must be installed on a filesystem client.

It includes:

- 'robinhood' daemon
- Reporting and other commands: 'rbh-report', 'rbh-find', 'rbh-du' and 'rbh-diff'
- configuration templates
- `/etc/init.d/robinhood` init script

Admin RPM (all purposes): `robinhood-adm`

Includes 'rbh-config' configuration helper. It is useful on the database host and Lustre MDS.

`robinhood-webgui` RPM installs a web interface to visualize stats from Robinhood database.

It must be installed on a HTTP server.

See section 8 for more details about this interface.

2.1.2 Build and install from the source tarball

It is advised to build a RPM from sources on your target system, so the program will have a better compatibility with your local Lustre and database version.

First, make sure the following packages are installed on your machine:

- `mysql-devel`
- lustre API library (if Robinhood is to be run on a Lustre filesystem):
'`/usr/include/liblustreapi.h`' and '`/usr/lib/liblustreapi.a`' are installed by Lustre rpm.

Retrieve Robinhood tarball from sourceforge: <http://sourceforge.net/projects/robinhood/files>

Unzip and untar the sources:

```
tar zxvf robinhood-2.5.2.tar.gz
cd robinhood-2.5.2
```

Then, use the “configure” script to generate Makefiles:

- use the `--with-purpose=TMPFS` option for using it as a temporary filesystem manager;

```
./configure --with-purpose=TMPFS
```

Other ‘./configure’ options:

- You can change the default prefix of installation path (default is /usr) using:
‘`--prefix=<path>`’
- If you want to disable Lustre specific features (getting stripe info, purge by OST...), use the ‘`--disable-lustre`’ option.

Finally, build the RPMs:

```
make rpm
```

RPMs are generated in the ‘`rpms/RPMS/<arch>`’ directory. RPM is eventually tagged with the lustre version it was built for.

NOTE: rpmbuild compatibility

Robinhood spec file (used for generating the RPM) is written for recent Linux distributions (RH5 and later). If you have troubles generating robinhood RPM (e.g. undefined rpm macros), you can switch to the older spec file (provided in the distribution tarball):

```
> mv robinhood.old_spec.in robinhood.spec.in
> ./configure ...
> make rpm
```

2.2 Robinhood service

Installing the rpm creates a ‘robinhood’ service. You can enable it like this:

```
> chkconfig robinhood on
```

This service starts one ‘robinhood’ instance for each configuration file it finds in ‘`/etc/robinhood.d/tmpfs`’ directory.

Thus, if you want to monitor several filesystems, create one configuration file for each of them. If there are common configuration blocks for those filesystems, you can use the ‘`%include`’ directive in configuration files.

NOTE: Suse Linux operating system

On SLES systems, the default dependency for boot scheduling is on “mysql” service.

However, in many cases, it could be too early for starting robinhood daemon, especially if the filesystem it manages is not yet mounted. In such case, you have to modify the following lines in `scripts/robinhood.init.sles.in` before you run `./configure`:

```
# Required-Start:    <required service>
```

2.3 Command line options

Legend:

- *> = new in robinhood 2.5
- ~> = changes in robinhood 2.5

Usage: robinhood [options]

Action switches:

- ~> **-S, --scan[=dir]**
Scan filesystem namespace.
If dir is specified, only scan the specified sub-directory.
- P, --purge**
Purge non-directory entries according to policy.
- C, --check-thresholds**
Only check thresholds of purge triggers without purging.
- R, --rmdir**
Remove directories according to policy.
- ~> **-r, --readlog[=mdt_idx] (Lustre 2.x only)**
Handle events from Lustre MDT ChangeLog.
If mdt_idx is specified, only read ChangeLogs for the given MDT.
Else, start 1 changelog reader thread per MDT (with DNE).

Default mode is: --scan --purge --rmdir

On Lustre 2:

Default mode is: --read-log --purge --rmdir

- *> **--diff=attrset** : when scanning or reading changelogs,
display changes for the given set of attributes (to stdout).
attrset is a list of options in:
path, posix, stripe, all, notimes, noatime.

Manual purge actions:

- purge-ost=ost_index,target_usage_pct**
Purge files on the OST specified by ost_index until it reaches the specified usage.
- purge-fs=target_usage_pct**
Purge files until the filesystem usage reaches the specified value.
- purge-class=fileclass**
Apply purge policies to files in the given class.

Behavior options:

- dry-run**
Only report actions that would be performed (rmdir, purge) without really doing them.
- i, --ignore-policies**
Force purging all eligible files, ignoring policy conditions.
- O, --once**
Perform only one pass of the specified action and exit.
- d, --detach**
Daemonize the process (detach from parent process).
- *> **--no-gc**
Garbage collection of entries in DB is a long operation when terminating a scan.
This option skips this operation if you don't care about removed entries (or don't expect entries to be removed).
This is also recommended for partial scanning (see **-scan=dir** option).

Config file options:

- f file, --config-file=file
Path to configuration file (or short name).
- T file, --template=file
Write a configuration file template to the specified file.
- D, --defaults
Display default configuration values.
- test-syntax
Check configuration file and exit.

Filesystem options:

- F path, --fs-path=path
Force the path of the filesystem to be managed (overrides configuration value).
- t type, --fs-type=type
Force the type of filesystem to be managed (overrides configuration value).

Log options:

- L logfile, --log-file=logfile
Force the path to the log file (overrides configuration value).
Special values "stdout" and "stderr" can be used.
- l level, --log-level=level
Force the log verbosity level (overrides configuration value).
Allowed values: CRIT, MAJOR, EVENT, VERB, DEBUG, FULL.

Miscellaneous options:

- h, --help
Display a short help about command line options.
- V, --version
Display version info
- p pidfile, --pid-file=pidfile
Pid file (used for service management).

Notes about the '-f' option:

- If no '-f' option is specified:
 - o robinhood first looks for RBH_CFG_DEFAULT environment variable.
 - o If this variable is not set, it searches for a config file in the directory '/etc/robinhood.d/tmpfs';
- you can specify a short name (e.g. 'foo') for a config file: it will first search for files named 'foo.conf' or 'foo.cfg' in '/etc/robinhood.d/tmpfs'. If none is found, it will search it in the current directory;
- You can specify a full path to the config file (eg. '-f /etc/robinhood.d/tmpfs/foo.conf')

2.4 Signals

Robinhood traps the following signals:

- **SIGTERM** (kill <pid>) and **SIGINT**: perform a clean shutdown;
- **SIGHUP** (kill -HUP <pid>): reload dynamic parameters from config file;
- **SIGUSR1** (kill -USR1 <pid>): dump process stats to its log file

2.5 Creating the database

Before running Robinhood for the first time, you must create its database.

- Install MySQL (mysql and mysql-server packages) on the node where you want to run the database engine.
- Start the database engine :
`service mysqld start`
- Use the 'rbh-config' command to check your configuration and create Robinhood database:

```
# check database requirements:
rbh-config precheck_db
```

```
# create the database:
rbh-config create_db
```

- If no option is given to `rbh-config`, it prompts for configuration parameters (interactive mode). Else, if you specify parameters on command line, it runs in batch mode.

Alternatively, you can perform the following steps of your own, without using the 'rbh-config' script:

- Create the database (one per filesystem) using the `mysqladmin` command:
`mysqladmin create <robinhood_db_name>`
- Connect to the database:
`mysql <robinhood_db_name>`

Then execute the following commands in the MySQL session:

- Create a robinhood user and set its password (MySQL 5+ only):
`create user robinhood identified by 'password';`
- Give access rights on database to this user (you can restrict client host access by replacing '%' by the node where robinhood will be running):
Mysql 5:
`GRANT USAGE ON robinhood_db_name.* TO 'robinhood'@'%';`
`GRANT ALL PRIVILEGES ON robinhood_db_name.* TO 'robinhood'@'%';`
Mysql 4.1:
`GRANT USAGE ON robinhood_db_name.* TO 'robinhood'@'%'`
`identified by 'password';`
`GRANT ALL PRIVILEGES ON robinhood_db_name.* TO 'robinhood'@'%';`
- The 'super' privilege is required for creating DB triggers (needed for accounting optimizations):
`GRANT SUPER ON *.* TO 'robinhood'@'%' IDENTIFIED BY 'password'`
`;`
- Refresh server access settings:
`FLUSH PRIVILEGES ;`

- o You can check user privileges using:
SHOW GRANTS FOR robinhood ;
- For testing access to database, execute the following command on the machine where robinhood will be running :
mysql --user=robinhood --password=password --host=db_host
robinhood_db_name
If the command is successful, a SQL shell is started. Else, you will get a 'permission denied' error.
- For now, the database schema is empty. Robinhood will automatically create it the first time it is launched.

2.6 Enabling Lustre Changelogs

With Lustre 2.x, file system scans are no longer required to update robinhood database. You just need to run an **initial scan** to populate robinhood DB. After that, robinhood gets events from Lustre using ChangeLog mechanism to update its DB. This avoids over-loading the filesystem with regular namespace scans!

If the MDS and the MGS are on the same host, you can simply enable this feature by running 'rbh-config' on this host:

```
> rbh-config enable_chglogs
```

Alternatively, this can be done by running the following commands:

On the MGS:

- Enable all changelog events (except ATIME, which is not to be used in production):
lctl conf_param mdd.*.changelog_mask all-ATIME

On the MDS:

- Changelogs consumers must be registered to Lustre to manage log records transactions properly. To do this, get a changelog reader id with the 'lctl' command:
>lctl
lctl > device lustre-MDT0000
lctl > changelog_register
lustre-MDT0000: Registered changelog userid 'c11'

Remember this id; it will be needed for writing PolicyEngine configuration file.

3 Writing configuration file

3.1 Syntax

General structure

The configuration file consists of several blocks. They can contain key/value pairs (separated by semi-colons), sub-blocks, Boolean expressions or set definitions (see '**set definitions**' below).

In some cases, blocks have an identifier.

```
BLOCK_1 bloc_id {
    Key = value;
    Key = value(opt1, opt2);
    Key = value;
    SUBBLOCK1 {
        Key=value;
    }
}
BLOCK_2 {
    (Key > value)
    and
    ( key == value or key != value )
}
CLASS_DEF {
    Set1 union Set2
}
```

Type of values

A value can be:

- A **string** delimited by single or double quotes (' or ").
- **[new 2.5]** An **environment variable**, starting with \$ (e.g. \$ROOT_PATH).
- A **Boolean** constant. Both of the following values are accepted and the case is not significant: TRUE, FALSE, YES, NO, 0, 1, ENABLED, DISABLED.
- A **numerical value** (decimal representation).
- A **duration**, i.e. a numerical value followed by one of those suffixes: 'w' for weeks, 'd' for days, 'h' for hours, 'min' for minutes, 's' for seconds. E.g.: 1s ; 1min ; 3h ; ... NB: if you do not specify a suffix, the duration is interpreted as seconds. E.g.: 60 will be interpreted at 60s, i.e.1 min.
- A **size**, i.e. a numerical value followed by one of those suffixes: PB for petabytes, TB for terabytes, GB for gigabytes, MB for megabytes, KB for kilobytes. No suffix is needed for bytes.
- A **percentage**: float value terminated by '%'. E.g.: 87.5%

Boolean expressions

Some blocks of configuration file are expected to be Boolean expressions on file attributes:

- AND, OR and NOT can be used in Boolean expressions.
- Brackets can be used for including sub-expressions.
- Conditions on attributes are specified with the following format:
<attribute> <comparator> <value>.
- Allowed comparators are '==', '<>' or '!=', '>', '>=', '<', '<='.

The following properties can be used in Boolean expressions:

- **tree**: entry is in the given filesystem tree. Shell-like wildcards are allowed.
Example:
tree == "/tmp/subdir/*/dir1" matches entry "/tmp/subdir/foo/dir1/dir2/foo" because

“/tmp/subdir/foo/dir1/dir2/foo” is in the “/tmp/subdir/foo/dir1” tree, that matches “/tmp/subdir/*/dir1” expression.

A tree can be specified using an absolute path (recommended) or a relative path (root path is the “fs_path” parameter in configuration).

The special wildcard “***” matches any count of directory levels:

E.g: tree == “**/.trash” matches any part of a filesystem under a “.trash” entry.

- **path:** entry exactly matches the path. Shell-like wildcards are allowed.
E.g: path == “/tmp/*/foo*” matches entry “/tmp/subdir/foo123”.

A path can be absolute (recommended) or relative (root is given by the “fs_path” parameter in configuration).

The special wildcard “***” matches any count of directory levels:

E.g: path == “**/.trash/**/file” matches any entry called “file” located somewhere under a “.trash” directory (at any depth).

- **name:** entry name matches the given regexp.
E.g: name == “*.log” matches entry “/tmp/dir/foo/abc.log”.
- **type:** entry has the given type (**directory**, **file**, **symlink**, **chr**, **blk**, **fifo** or **sock**).
E.g: type == “symlink”.
- **owner:** entry has the given owner (owner name expected).
E.g: owner == “root”.
- **group:** entry is owned by the given group (group name expected).
- **size:** entry has the specified size. Value can be suffixed with KB, MB, GB...
E.g: size >= 100MB matches file whose size equals 100x1024x1024 bytes or more.
- **last_access:** condition based on the last access time to a file (for reading or writing). This is the difference between current time and max(ctime, mtime, atime). Value can be suffixed by ‘sec’, ‘min’, ‘hour’, ‘day’, ‘week’...
E.g: last_access < 1h matches files that have been read or written within the last hour.
- **last_mod:** condition based on the last modification time to a file. This is the difference between current time and max(ctime, mtime).
E.g: last_mod > 1d matches files that have not been modified for more than a day.
- **creation:** condition based on file creation time. With scanning mode, this value is an estimation based on the time when robinhood sees a file for the first time and ctime.
E.g.: creation > 1h matches files created more than 1 hour ago.
- **last_archive** (HSM and backup modes only): condition on the last time the file has been archived. The special value of 0 matches files that have never been archived.
E.g. last_archive > 1d matches files that have been archived more than 1 day ago.

- **last_restore** (HSM mode only): condition on the last time the file has been restore from the backend storage. E.g. `last_restore < 1d` matches files that have been restored in the last 24 hours.
- **ost_index**: condition on OSTs where a file is stored. A file can be striped on several OSTs, thus:
`"ost_index == N"` is true if at least one part of the file is stored in OST index #N.
`"ost_index != N"` is true if no part of the file is stored in OST index #N.
 (Note: `<` and `>` are not allowed for this criteria).
- **ost_pool**: condition about the OST pool name where the file was created. Wildcarded expressions are allowed.
 E.g. `ost_pool == "pool*"`.
- **xattr.xxx**: test the value of a user-defined extended attribute of the file.
 E.g. `xattr.user.tag_no_purge == "1"`
 - xattr values are interpreted as text string;
 - regular expressions can be used to match xattr values;
 E.g. `xattr.user.foo == "abc.[1-5].*"` matches file having xattr `user.foo = "abc.2.xyz"`
 - if an extended attribute is not set for a file, it matches empty string.
 E.g. `xattr.user.foo == ""` ⇔ xattr 'user.foo' is not defined
- **dircount** (for directories only): the directory has the specified number of entries (except `'.'` and `'..'`).
 E.g. `dircount > 10000` matches directories with more that 10 thousand child entries.

Example of Boolean expression:

```
ignore {
  ( name == "*.log" and size < 15GB )
  or ( owner == "root" and last_access < 2d )
  or not tree == "/fs/dir"
}
```

Set definitions

In the case of FileClass definitions, you can define FileClasses as the union or intersection of other FileClasses previously defined. This can be done using “**union**”, “**inter**” and “**not**” keywords. Such expressions can be encapsulated between parenthesis.

Example:

```
FileClass my_set_union {
  definition { ( Class1 union Class2 ) inter ( not Class3 ) }
}
```

Comments

The '#' and '/' signs indicate the beginning of a comment (except if there are in a quoted string). The comment ends at the end of the line.

E.g.:

```
# this is only a comment line
```

```
x = 32 ; # a comment can also be placed after a significant line
```

Includes

A configuration file can be included from another file using the '%include' directive. Both relative and absolute paths can be used.

E.g.:

```
%include "subdir/common.conf"
```

Configuration blocks

The main blocks in a configuration file are:

- **General** (mandatory): main parameters.
- **Log**: log and alert parameters (log files, log level...).
- **Filesets**: definition of file classes
- **Purge_Policies**: defines purge policies.
- **Purge_Trigger**: specifies conditions for starting purges.
- **Purge_Parameters**: general options for purge.
- **Rmdir_Policy**: defines empty directory removal policy.
- **Rmdir_Parameters**: options about empty directory removal.
- **ListManager** (mandatory): database access configuration.
- **FS_Scan**: options about scanning the filesystem.
- **db_update_policy**: parameters about file class periodic matching, and entry information update interval.
- **ChangeLog**: parameters related to Lustre 2.x changelogs.
- **EntryProcessor**: configuration of entry processing pipeline (for FS scan).

Those blocks are described in the following sections.

3.2 Configuration template and default parameters

Template file

To easily create a configuration file, you can generate a documented template using the --template option of robinhood, and edit this file to set the values for your system:

```
robinhood --template=<template file>
```

Default configuration values

To know the default values of configuration parameters use the --defaults option:

```
robinhood --defaults
```

3.3 General parameters

General parameters are set in a configuration block whose name is '**General**'.

The following parameters can be specified in this block:

- **fs_path** (string, mandatory): the path of the file system to be managed. This must be an absolute path. This parameter can be overridden by "--fs-path" parameter on command line.
E.g.: `fs_path = "/tmp_fs";`
- **fs_type** (string, mandatory): the type of the filesystem to be managed (as displayed by mount). This is mainly used for checking if the filesystem is mounted. This parameter can be overridden by "--fs-type" parameter on command line.
E.g.: `fs_type = "lustre";`
- **fs_key**: this indicates the filesystem property used as unique and persistent file system identifier. Possible values are: 'fsname', 'devid' or 'fsid' ('fsid' is NOT recommended as it may change at each *mount*).
E.g.: `fs_key = fsname;`
- **stay_in_fs** (Boolean): if this parameter is TRUE, robinhood checks that the entries it handles are in the same device as *fs_path*, which prevents from traversing mount points.
E.g.: `stay_in_fs = TRUE;`
- **check_mounted** (Boolean): if this parameter is TRUE, robinhood checks that the filesystem of *fs_path* is mounted.
E.g.: `check_mounted = TRUE;`
- **lock_file** (string): robinhood suspends its activity when this file exists.
E.g.: `lock_file = "/var/lock/robinhood.lock";`

3.4 Log and alerts parameters

The parameters described in this section can be specified in the '**Log**' block. They control logging, reporting and alerts.

Logging parameters

- **debug_level** (string): verbosity level of logs. This parameter can be overridden by "--log-level" parameter on command line.
Allowed values are :
 - FULL: highest level of verbosity. Trace everything.
 - DEBUG: trace information for debugging.
 - VERB: high level of traces (but usable in production).
 - EVENT: standard production log level.
 - MAJOR: only trace major events.
 - CRIT: only trace critical events.

E.g.: `debug_level = VERB;`

- **log_file** (string): file where logs are written. This parameter can be overridden by “--log-file” parameter on command line.

E.g.: `log_file = "/var/logs/robinhood/robinhood.log";`

- **report_file** (string): file where purge and rmdir operations are logged.

E.g.: `report_file = "/var/logs/robinhood/purge_report.log";`

Notes about **log_file** and **report_file**:

- ⇒ Make sure the log directory exists.
- ⇒ robinhood is compliant with log rotation (if its log file is renamed, it will automatically open a new file).
- ⇒ The following special values can be used as log files:
 - o **'stdout'**: log to standard output
 - o **'stderr'**: log to standard error
 - o **'syslog'**: log using syslog
- ⇒ For syslog, you can select the syslog facility using the **'syslog_facility'** parameter.
E.g.:
`log_file = syslog ;`
`syslog_facility = local1.info ;`

[new 2.5.1] Log header control

In the previous versions of robinhood, log lines always had the following format:

`<date> <time> <process name>@<hostname>[pid/thrd]: <module> | <log msg>`

Now the default log header is shorter:

`<date> <time> [pid/thrd] <module> | <log msg>`

but you can still add additional info in the log using new configuration parameters:

- **log_procname** (boolean): display the process name in the log header (default : 'no').
- **log_hostname** (boolean): display the host name in the log header (default : 'no').
- **log_module** (boolean): display the module name in the log header (default : 'yes').

Alert parameters

Two methods can be used for raising alerts: sending a mail, writing to a file, or both.

This is set by the following parameters:

- **alert_file** (string): if this parameter is set, alerts are written to the specified file.
E.g.: `alert_file = "/var/logs/robinhood/alerts.log";`
- **alert_mail** (string): if this parameter is set, mail alerts are sent to the specified recipient.
E.g.: `alert_mail = "admin@localdomain";`

Alert parameters:

- **alert_show_attrs** (Boolean): If true, details entry attributes in alerts.
- **batch_alert_max** (integer): this controls alert batching (sending 1 alert summary instead of 1 per entry):
 - If the value is 0, there is no limit in batching alerts. 1 summary is send after each scan.
 - If the value is 1, alerts are not batched.
 - If the value is $N > 1$, a summary is sent every N alerts.

3.5 File classes

You may need to apply different purge policies depending on file properties. To do this, you can define file classes.

A file class is defined by a 'FileClass' block. All file class definitions must be grouped in the 'Filesets' block of the configuration file.

Each file class has an identifier (that you can use for addressing it in policies) and a definition (a condition for entries to be in this file class).

FileClasses can be defined as the union or the intersection of other FileClasses, using 'inter' and 'union' keywords in fileclass definition.

File classes definition overview:

```
Filesets {
    FileClass my_class_1 {
        Definition {
            tree == "/fs/dir_A"
            and
            owner == root
        }
    }

    FileClass my_class_2 {
        ...
    }

    FileClass my_inter_class {
        Definition { my_class_1 inter my_class_3 }
    }
    ...
}
```

Important note: if you modify fileclass definitions or target fileclasses of policies, you need to reset fileclass information in Robinhood database.

To do so, run the following command:

```
rbh-config reset_classes
```

3.6 Purge policies

In general, files are purged in the order of their last access time (LRU list). You can however specify conditions to allow/avoid entries to be purged, depending on their file class, and file properties.

To define purge policies, you can specify:

- Sets of entries that must never be purged (ignored).
- Purge policies to be applied to file classes.
- A default purge policy for entries that don't match any file class.

In configuration file, all those parameters are grouped in a '**Purge_Policies**' block that consists in:

- '**Ignore**' sub-blocks: Boolean expressions to "white-list" filesystem entries depending on their properties.

E.g.: `Ignore { size == 0 or type == "symlink" }`

- **‘Ignore_fileclass’**: “white-list” all entries of a fileclass (see section 3.5 about defining File classes).

E.g.: `Ignore_FileClass = my_class_1;`

- **‘Policy’** sub-blocks: specify conditions for purging entries of file classes. A policy has a custom name, one or several target file classes, and a condition for purging files.

E.g:

```
Policy purge_classes_2and3
{
    target_fileclass = class_2;
    target_fileclass = class_3;

    condition
    {
        Last_access > 1h
    }
}
```

- A default policy that applies to files that don’t match any previous file class or ‘ignore’ directive. It is a special ‘Policy’ block whose name is ‘default’ and with no target_fileclass.

E.g:

```
Policy default
{
    condition
    {
        last_access > 30min
    }
}
```

As a summary, the ‘purge_policies’ block will look like this:

```
purge_policies
{
    # don't purge symlinks and entries owned by root
    Ignore { owner == "root" or type == symlink }

    # don't purge files of classes 'class_xxx' and 'class_yyy'
    Ignore_FileClass = class_xxx ;
    Ignore_FileClass = class_yyy ;

    # purge policy for files of 'my_class1' and 'my_class2'
    policy my_purge_policy1
    {
        target_fileclass = my_class1;
        target_fileclass = my_class2;
        condition { last_access > 1h and last_mod > 2h }
    }
    ...
    # purge policy for other files
    policy default
    {
        condition { last_access > 10min }
    }
}
```

Note: the target fileclasses are matched in the order they appear in the purge_policies block, so make sure to specify the more restrictive classes first.

Thus, in the following example, the second policy can’t never be matched, because A already matches all entries in A_inter_B:

```

Filesets {
    Fileclass A { ... }
    Fileclass B { ... }
    Fileclass A_inter_B { definition { A inter B } }
}

purge_policies
{
    policy purge_1
    {
        target_fileclass = A; # all entries of fileclass A
        ...
    }
    policy purge_2
    {
        target_fileclass = A_inter_B; # never matched!!!
        ...
    }
}

```

3.7 *Purge triggers*

Triggers describe conditions for starting/stopping purges. They are defined by ‘purge_trigger’ blocks. Each trigger consists of:

- The type of condition (on global filesystem usage, on OST usage, on volume used by a user or a group...);
- A purge start condition ;
- A purge target condition ;
- An interval for checking start condition.
- Notification options

Several triggers can be specified.

Type of condition

The type of condition is specified by “**trigger_on**” parameter.

Possible values are:

- **global_usage:** purge start/stop condition is based on the space used in the whole filesystem (based on *df* return). All entries in filesystem are considered for such a purge.
- **OST_usage:** purge start/stop condition is based on the space used on each OST (based on *lfs df*). Only files stored in an OST are considered for such a purge.
- **user_usage[user1, user2...]:** purge start/stop condition is based on the space used by a user (kind of quota). Only files that are owned by a user over the limit are considered for such a purge. If it is used with no arguments, all users will be affected by this policy.
A list of users can also be specified for restricting the policy to a given set of users (comma-separated list of users between brackets). [Fully implemented since Robinhood 2.2].
- **group_usage[grp1, grp2...]:** purge start/stop condition is based on the space used by a group (kind of quota). Only files that are owned by a group over the limit are considered for purge. If it is used with no arguments, all groups will be affected by this policy.

A list of groups can also be specified for restricting the policy to a given set of groups (coma-separated list of groups between brackets). [Fully implemented since Robinhood 2.2].

- **periodic:** purge runs at scheduled interval, with no condition on filesystem usage.

Start condition

This is mandatory for all types of conditions.

A purge start condition can be specified by two ways: percentage or volume.

- **high_threshold_pct** (percentage): specifies a percentage of space used over which a purge is launched.
- **high_threshold_vol** (size): specifies a volume of space used over which a purge is launched. The value for this parameter can be suffixed by KB, MB, GB, TB...
- **high_threshold_cnt** (count): condition based on the number of inodes in the filesystem. It can be used for **global_usage**, **user_usage** and **group_usage** triggers. The value can be suffixed by K, M, ...

No threshold is expected for 'periodic' triggers.

Stop condition

This is mandatory for all types of conditions.

A purge stop condition can also be specified by two ways: percentage or volume.

- **low_threshold_pct:** specifies a percentage of space used under which a purge stops.
- **low_threshold_vol:** specifies a volume of space used under which a purge stops. The value for this parameter can be suffixed by KB, MB, TB... (the value is interpreted as bytes if no suffix is specified).
- **low_threshold_cnt** (count): condition based on the number of inodes in the filesystem. It can be used for **global_usage**, **user_usage** and **group_usage** triggers. The value can be suffixed by K, M, ...

No threshold is expected for 'periodic' triggers.

Runtime interval

The time interval for checking a condition is set by parameter "**check_interval**". The value for this parameter can be suffixed by 'sec', 'min', 'hour', 'day', 'week', 'year'... (the value is interpreted as seconds if no suffix is specified).

Raise alert when high threshold is reached

Optionally, an alert can be raised each time the high threshold is reached. This can be done by setting the "**alert_high**" parameter (Boolean) in the trigger:

```
alert_high = TRUE;
```

Don't raise alert when low threshold cannot be reached

By default, robinhood raises an alert if it can't purge enough data to reach the low threshold.

You can disable these alerts by adding this in a trigger definition:

```
alert_low = FALSE ;
```

Examples

Check ‘df’ every 5 minutes, start a purge if space used > 85% of filesystem and stop purging when space used reaches 84.5%:

```
Purge_Trigger
{
    trigger_on = global_usage ;
    high_threshold_pct = 85% ;
    low_threshold_pct = 84.5% ;
    check_interval = 5min ;
}
```

Check OST usage every 5 minutes, start a purge of files on an OST if it space used is over 90% and stop purging when space used on the OST falls to 85%:

```
Purge_Trigger
{
    trigger_on = OST_usage ;
    high_threshold_pct = 90% ;
    low_threshold_pct = 85% ;
    check_interval = 5min ;
}
```

Daily check the space used by each user of a given list. If one of them uses more than 1TB, its files are purged until it uses less than 800GB. Also send an alert in this case.

```
Purge_Trigger
{
    trigger_on = user_usage(foo, charlie, roger, project*) ;
    high_threshold_vol = 1TB ;
    low_threshold_vol = 800GB ;
    check_interval = 1day ;
    alert_high = TRUE ;
}
```

Check that user inode usage is less than 100k entries (and send a notification in this case):

```
Purge_Trigger
{
    trigger_on = user_usage ;
    high_threshold_cnt = 100k ;
    low_threshold_cnt = 100k ;
    check_interval = 1day ;
    alert_high = TRUE ;
}
```

Apply purge policies twice a day (whatever the filesystem usage):

```
Purge_Trigger
{
    trigger_on = periodic;
    check_interval = 12h;
}
```

Note: check triggers conditions without purging

If robinhood is started with the ‘--check-thresholds’ option instead of ‘--purge’, it will only check for trigger conditions and eventually send notifications, without purging data.

3.8 Purge parameters

Purge parameters are specified in a ‘purge_parameters’ block.

The following options can be set:

- **nb_threads_purge** (integer): this determines the number of purge operations that can be performed in parallel.
E.g.: `nb_threads_purge = 8 ;`
- **post_purge_df_latency** (duration): immediately after purging data, *df* and *ost df* may return a wrong value, especially if freeing disk space is asynchronous. So, it is necessary to wait for a while before issuing a new *df* or *ost df* command after a purge. This duration is set by this parameter.
E.g.: `post_purge_df_latency = 1min ;`
- **purge_queue_size** (integer): this advanced parameter is for leveraging purge thread load.
- **db_result_size_max** (integer): this impacts memory usage of MySQL server and Robinhood daemon. The higher it is, the more memory they will use, but less DB requests will be needed.
- **recheck_ignored_classes** (Boolean): this results in checking previously ignored entries when applying purge policies (default behavior). Disable it to speed-up policy application.

3.9 Specifying 'rmdir' policy

Directory removal is driven by the 'rmdir_policy' section in the configuration file:

- **age_rm_empty_dirs** (duration): indicates the time after which an empty directory is removed. If set to 0, empty directory removal is disabled.
- You can specify one or several '**ignore**' condition for directories you never want to be removed.
- '**recursive_rmdir**' sub-blocks indicates that the matching directories must be removed recursively. /!\In this case, the whole directories content is removed without checking policies on their content (whitelist rules...).

Example:

```
rmdir_policy
{
    # remove empty directories after 15 days
    age_rm_empty_dirs = 15d;

    # recursively remove ".trash" directories after 15 days
    recursive_rmdir
    {
        name == ".trash" and last_mod > 15d
    }

    # whitelist directories matching the following condition
    ignore
    {
        depth < 2
        or
        owner == 'foo'
        or
        tree == /fs/subdir/A
    }
}
```

```
}
}
```

3.10 'Rmdir' parameters

Directory removal parameters are specified in the 'rmdir_parameters' block.

The following options can be set:

- **runtime_interval** (duration): interval for performing empty directory removal.
- **nb_threads_rmdir** (integer): this determines the number of 'rmdir' operations that can be performed in parallel.
E.g.: `nb_threads_rmdir = 4;`
- **rmdir_op_timeout** (duration): this specifies the timeout for 'rmdir' operations. If a thread is stuck in a filesystem operation during this time, it is cancelled.
E.g.: `rmdir_op_timeout = 15min;`
- **rmdir_queue_size** (integer): this advanced parameter is for leveraging rmdir thread load.

3.11 Periodic fileclass matching

By default, Robinhood matches fileclasses for an entry each time it updates its attributes. Alternatively you can match fileclasses only if it has not been updated for a given time.

The fileclass matching interval is set using the 'fileclass_update' parameter in the new 'db_update_policy' section:

```
db_update_policy
{
    fileclass_update = periodic( 1h );
}
```

Possible values are:

- `never`: match file class once, and never again
- `always`: always re-match file class when applying policies (this is the old behavior)
- `periodic(<period>)`: periodically re-match fileclasses

3.12 Periodic information update when handling Lustre changelogs

This is similar to fileclass periodic matching, but it is for updating file metadata and path in the database when processing Lustre changelogs.

Indeed, if robinhood receives a lot of events for a given entry, it can be very loud to update entry information when processing each event.

You can specify the way entry path and entry metadata is updated in the database using 'md_update' and 'path_update' parameters in the new 'db_update_policy' section:

```
db_update_policy
{
    md_update = on_event_periodic(1sec,1min);
```



```

    path_update = on_event;
}

```

Possible values for those parameters are:

- **never**: retrieve it once, then never update the information
- **always**: always update the information when receiving an event for the entry
- **periodic(<period>)**: update the information only if it was not updated for a while
- **on_event**: update the information every time the event is related to it (eg. update entry path when the event is a 'rename' on the entry).
- **on_event_periodic(<interval_min>, <interval_max>)**: this is the smarter one

3.13 Database parameters

The 'ListManager' block is the configuration for accessing the database.

ListManager parameters:

- **commit_behavior**: this is the method for committing information to database.
The following values are allowed:
 - **transaction**: group operations in transactions (best consistency, recommended).
 - **autocommit**: more efficient with some DB engines, but database inconsistencies may appear.
 - **periodic(<nbr_transactions>)**: operations are packed in large transactions before they are committed. 'Commit' is done every n transactions. This method is more efficient for in-file databases like SQLite. This causes no database inconsistency, but more operations are lost in case of a crash.
E.g: `commit_behavior = periodic(1000);`
- **connect_retry_interval_min, connect_retry_interval_max** (durations):
'connect_retry_interval_min' is the time (in seconds) to wait before re-establishing a lost connection to database. If reconnection fails, this time is doubled at each retry, until 'connect_retry_interval_max'.
E.g: `connect_retry_interval_min = 1;`
 `connect_retry_interval_max = 30;`

Accounting parameters (in ListManager):

- **user_acct, group_acct** (Booleans): these parameters enable or disable optimized reports for user and group statistics.
By default these parameters are enabled. If you disable them (or one of them), report generation will be slower, but it will make database operation faster during filesystem scans. For instance, if you only need reports by user, disable *group_acct* to optimize scan speed.
E.g: `user_acct = on ;`
 `group_acct = off ;`
See section 10.3 for more details.

MySQL specific configuration is set in a 'MySQL' sub-block, with the following parameters:

- **server**: machine where MySQL server is running. Both server name and IP address can be specified.
E.g.: `server = "mydbhost.localnetwork.net";`
- **db** (string, mandatory): name of the database.
E.g.: `db = "robinhood_db";`
- **user** (string): name of the database user.
E.g.: `user = "robinhood";`
- **password** or **password_file** (string, mandatory): there are two methods for specifying the password for connecting to the database, depending of the security level you want. You can directly write it in the configuration file, by setting the '**password**' parameter. You can also write the password in a distinct file (with more restrictive rights) and give the path to this file by setting '**password_file**' parameter. This makes it possible to have different access rights for config file and password file.
E.g.: `password_file = "/etc/robinhood/.dbpass";`
- **innodb** (Boolean): by default, Robinhood creates a database using InnoDB MySQL engine. Usually, InnoDB results in better performances, depending on your usage of robinhood.
If you want to use MyISAM instead, disable this parameter:
`innodb = disabled ;`
Note: if Robinhood already created its database tables, you need to drop them (run `rbh-config empty_db`) and restart Robinhood to create the tables with the new DB engine.

3.14 Filesystem scan parameters

Parameters for scanning the filesystem are set in the '**FS_Scan**' block.
It can contain the following parameters:

- **scan_interval** (duration): specifies a fix frequency for scanning the filesystem (daemon mode).
- or
- **min_scan_interval**, **max_scan_interval** (durations): it is possible to adapt the scan frequency depending on the current filesystem usage. Indeed, it is not necessary to scan the filesystem frequently when it is empty (because no purge is needed). When the filesystem is full, robinhood will need a fresh list for purging files, so it is better to scan more frequently. For this, specify the interval between scans using:
 - **min_scan_interval**: the frequency for scanning when filesystem is full;
 - **max_scan_interval**: the frequency for scanning when filesystem is empty
 The interval between scans is computed according to this formula:

$$\text{min} + (100\% - \text{current usage}) * (\text{max} - \text{min})$$
 Note : those parameters are not compatible with simple **scan_interval**.
 - **nb_threads_scan** (integer): number of threads used for scanning the filesystem in parallel.
 - **scan_retry_delay** (duration): if a scan fails, this is the delay before starting another.
 - **scan_op_timeout** (duration): this specifies the timeout for `readdir/getattr` operations. If filesystem operations are blocked more than this time, there are cancelled. It is recommended to enable '`exit_on_timeout`' option in that case.

- **exit_on_timeout** (Boolean): robinhood exits if filesystem operations are blocked for a long time (specified by 'scan_op_timeout').
- **spooler_check_interval** (duration): interval for testing FS scans, deadlines and hangs.
- **nb_prealloc_tasks** (integer): number of pre-allocated task structures (advanced parameter).
- **completion_command** (string): external completion command to be called when robinhood terminate a filesystem scan. The full path for the command must be specified. Command arguments can contain special values:
 - o {cfg} for the config file
 - o {fspath} for the path of the filesystem managed by robinhood
 Eg.: completion_command = "/path/to/script.sh -f {cfg} -p {fspath}";
- **Ignore** block (Boolean expression): robinhood will skip entries and directories that match the given expression. Several ignore blocks can be defined in FS_Scan.

Examples:

```
Ignore {
    # ignore ".snapshot" directories (don't scan them)
    type == directory
    and
    name == ".snapshot"
}

Ignore {
    # ignore a whole part of the filesystem
    tree == "/mnt/lustre/dont_scan_me"
}
```

3.15 Lustre 2.x changelog parameters

With Lustre 2.x, FS scan are no longer required to update robinhood's database. Reading Lustre's changelog is much more efficient, because this does not load the filesystem as much as a full namespace scan. However, **an initial scan is still required** to initially populate robinhood DB.

Accessing the chngelog is driven by the '**ChangeLog**' block of configuration.

It contains one 'MDT' block for each MDT, with the following information:

- **mdt_name** (string): name of the MDT to read ChangeLog from (basically "MDT0000").
- **reader_id** (string): log reader identifier, returned by '**lctl changelog_register**' (see section 2.6: "Enabling Lustre Changelogs").

So far, changelog readers need to perform active polling to get new events from MDT. Until that changes, you need to activate polling:

force_polling = ON;

The polling interval control the frequency for reading changelog records from the MDTs:

polling_interval = 1s;

You can also control the frequency of acknowledging changelog records to Lustre (this reduces the number of filesystem calls), by specifying “**batch_ack_count**”.

A zero value indicates that changelog records are acknowledged once they have all been read and processed.

E.g: clear changelog every 500 records

```
batch_ack_count = 500 ;
```

[new 2.5.3] Robinhood can log every changelog record it gets from Lustre. This can be useful for debugging, or to analyze filesystem access patterns. To enable this, set the following parameter: `dump_file = "/path/to/dump_file";`

A basic “ChangeLog” block looks like this:

```
ChangeLog {
    MDT {
        mdt_name = "MDT0000";
        reader_id = "cl1";
    }
    # In the case of using DNE, you can define other MDTs
    # as other MDT blocks.
    MDT {
        mdt_name = "MDT0001";
        reader_id = "cl2";
    }
}
```

[new 2.5] Lustre Changelog batching (optimization)

To speed up changelog processing, robinhood retains changelog records in memory a short time, to aggregate similar/redundant Changelog records on the same entry before updating its database (e.g. both MTIME and CLOSE events mean the file mtime and size may have changed and must be updated).

You can configure this batching using the following parameters in the ‘Changelog’ block:

- **queue_max_size** (integer): the maximum number of retained records. Default is 1000.
- **queue_max_age** (duration): the maximum time a record is retained. Default is 5s.
- **queue_check_interval** (duration): the period for checking records age (related to queue_max_age parameter). Default is 1s.

[new 2.5] Reading Changelogs with Lustre DNE

For handling several Lustre MDTs (with DNE), there are 2 possible configurations:

- You can run Changelog readers for all MDTs as a single Robinhood process. In this case, robinhood will have 1 changelog reader thread per MDT in the same Robinhood process. To do so, simply run Robinhood as usual (with no option, or with --readlog option).
- You can run 1 changelog reader process per MDT, possibly distributed on different hosts. To do so, give the MDT index as an option to --readlog.

Example:

on host1, to read changelogs from the first MDT in config file (should be MDT0000):

```
robinhood --readlog=0
```

on host2, to read changelogs from the second MDT (should be MDT0001):
robinhood --readlog=1
and so on...

[new 2.5] Dealing with old Changelogs bugs

In some past Lustre versions, there were a couple of lacks or defects in Changelog records. If your MDS runs an old version of Lustre, you can specify whether your MDS has the patches for the following bugs: LU-543, LU-1331.

This can be specified using the following parameters in the ‘Changelog’ block:

- **mds_has_lu543** (boolean): enable if the MDS has the fix for LU-543. Disable if it has the bug.
- **mds_has_lu1331** (boolean): enable if the MDS has the fix for LU-1331. Disable if it has the bug.

If you don’t know, don’t specify those parameters. Robinhood will guess them by taking a look at record contents. In this case it will display such messages in the log:

Detected LU-543.

which means the bug is not present on your MDS.

3.16 Entry processor pipeline options

When scanning a filesystem or reading changelogs, robinhood process incoming information using a pool of worker threads. The processing is based on a pipeline model, with 1 stage for each kind of operation (FS operation, DB operation ...).


The behavior of this module is controlled by the ‘**EntryProcessor**’ block.

[new 2.5] You have the choice between 2 strategies to maximize robinhood processing speed:


- *multi-threading*: perform multiple DB operations in parallel as independent transactions.
- *batching*: batch database operations (insert, update...) into a single transaction, which minimize the need for IOPS on the database backend. Batches are not executed in parallel.

The following benchmark evaluated the DB performance for each strategy. This benchmark ran on a simple test-bed using basic SATA disk as DB storage for innodb.

DB operations multithreading



DB operations batching



The results show that batching is more efficient than multi-threading whatever the thread count, so it has been made the default behavior for robinhood 2.5.

You can control batches size using this parameter:

- **[new 2.5] max_batch_size** (positive integer): by default, the entry processor tries to batch similar database operations to speed them. This can be controlled by the `max_batch_size` parameter. The default max batch size is 1000.

If your DB storage backend is efficient enough (high IOPS) you may want to give a try to the other strategy. To switch from *batching* strategy to *multi-threading*, set **max_batch_size = 1**. This will automatically disable batching and enables multi-threading for DB operations. Consider increasing **nb_threads** parameter in this case (see below).

You can control the number of worker thread for the pipeline using the following parameter:

- **nb_threads** (integer): total number of threads for performing pipeline tasks. Default is 4. Consider increasing it if you disable *batching*.

You can also precisely control the max number of threads that work for each pipeline step, using those parameters:

- **STAGE_GET_FID_threads_max** (integer): when scanning a filesystem, this indicates the maximum number of threads that can perform a `llapi_path2fid` operation simultaneously.
- **STAGE_GET_INFO_DB_threads_max** (integer): this limits the number of threads that simultaneously check if an entry already exists in database.
- **STAGE_GET_INFO_FS_threads_max** (integer): this limits the number of threads that simultaneously retrieve information from filesystem (`getstripe...`).
- **STAGE_REPORTING_threads_max** (integer): this limits the number of threads that simultaneously check and raise alerts about filesystem entries.
- **STAGE_PRE_APPLY_threads_max** (integer): this step is only filtering entries before the `DB_APPLY` step. No reason to limit it.
- **STAGE_DB_APPLY_threads_max** (integer): this limits the number of threads that simultaneously insert/update entries in the database.

E.g.: for limiting the number of simultaneous database operations:
`STAGE_DB_APPLY_threads_max = 4;`

Other parameters:

- **max_pending_operations** (integer): this parameter limits the number of pending operations in the pipeline, so this prevents from using too much memory. When the number of queued entries reaches this value, the scanning process is slowed-down to keep the pending operation count below this value.

Alerts

One of the tasks of the Entry Processor is to check alert rules and raise alerts. For defining an alert, simply write an **Alert** sub-block with a Boolean expression that describes the condition for raising an alert (see section 3.1 for more details about writing Boolean expressions on file attributes).

Alerts can be named, to easily identify/distinguish them in alert summaries.

E.g.: raise an alert if a directory contains more than 10 thousand entries:

```
Alert    Large_flat_directory
{
    type == directory
    and
    dircount > 10000
    and
    last_mod < 2h
}
```

Another example: raise an alert if a file is larger than 100GB (except for user 'foo'):

```
Alert    Big_File
{
    type == file
    and
    size > 100GB
    and
    owner != 'foo'
    and
    last_mod < 2h
}
```

Tip: alert rules are matched at every scan. If you don't want to be alerted about a given file at each scan, it is advised to specify a condition on `last_mod`, so you will only be alerted for recently modified entries.

FileClass matching

By default, entry fileclass are matched immediately when entries are discovered (basically at scan time). This makes it possible for the reporting command to display the fileclass information.

If the filesystem is overloaded, this can be disabled using the **'match_classes'** parameter in the **'EntryProcessor'** block. In this case, entry fileclasses will only be matched when applying policies.

```
match_classes = FALSE;
```

Detecting “fake” mtime

You may have robinhood policies based on file modification time. If users change their file modification times using 'touch' command, 'utime' call, or copying files using 'rsync' or 'cp -p', it may result in unexpected policy decisions (like purging recently modified files because the user set a mtime in the past).

To avoid such situation, enable the **'detect_fake_mtime'** parameter in the **'EntryProcessor'** block:

```
detect_fake_mtime = TRUE;
```

Note: robinhood traces the fake mtime it detects with 'debug' trace level.

4 Reporting tool (rbh-report)

4.1 Overview

The content of Robinhood's database can be very useful for building detailed reports about filesystem content. For example, you can know how many entries of each type (directory, file, symlink...) exist in the filesystem, the min/max/average size of files, the min/max/average count of entries in directories, the space used by a given user, etc...

All those statistics can easily be retrieved using Robinhood reporting tool: **rbh-report**.

4.2 Command line

Legend:

*> = new in robinhood 2.5

~> = changes in robinhood 2.5

Usage: rbh-report [options]

Stats switches:

- activity, -a**
Display stats about daemon's activity.
- fs-info, -i**
Display filesystem content statistics.
- *> **--entry-info path|id, -e path|id**
Display all stored information about an entry.
- class-info [=fileclass]**
Display Fileclasses summary. Use optional parameter fileclass for retrieving stats about a given fileclass (wildcards allowed).
- user-info [=user], -u user**
Display user statistics. Use optional parameter user for retrieving stats about a single user.
- group-info [=group], -g group**
Display group statistics. Use optional parameter group for retrieving stats about a single group.
- top-dirs [=count], -d count**
Display largest directories. Optional argument indicates the number of directories to be returned (default: 20).
- top-size [=count], -s count**
Display largest files. Optional argument indicates the number of files to be returned (default: 20).
- top-purge [=count], -p count**
Display oldest entries eligible for purge. Optional argument indicates the number of entries to be returned (default: 20).
- top-rmdir [=count], -r count**
Display oldest empty directories eligible for rmdir. Optional argument indicates the number of dirs to be returned (default: 20).
- top-users [=count], -U count**
Display top disk space consumers. Optional argument indicates the number of users to be returned (default: 20).
- dump, -D**
Dump all filesystem entries.
- dump-user user**

- Dump all entries for the given user.
- dump-group group**
 Dump all entries for the given group.
- ~> **--dump-ost ost_index|ost_set**
 Dump all entries on the given OST or set of OSTs (e.g. 3,5-8).

Filter options:

The following filters can be specified for reports:

- P path, --filter-path path**
 Display the report only for objects in the given path.
- C class, --filter-class class**
 Report only entries in the given FileClass.
- count-min cnt**
 Display only topuser/userinfo with at least cnt entries

Accounting report options:

- size-profile, --szprof**
 Display size profile statistics
- by-size-ratio range, --by-szratio range**
 Sort on the ratio of files in the given size-range
 range: <val><sep><val>- or <val><sep><val-1> or <val><sep>inf
 <val>: 0, 1, 32, 1K 32K, 1M, 32M, 1G, 32G, 1T
 <sep>: ~ or ..
 e.g: 1G..inf, 1..1K-, 0..31M
- by-count**
 Sort top users by count instead of sorting by volume
- by-avgsz**
 Sort users by average file size
- reverse**
 Reverse sort order
- S, --split-user-groups**
 Display the report by user AND group
- F, --force-no-acct**
 Generate the report without using accounting table

Config file options:

- config-file=file, -f file**
 Path to configuration file (or short name).

Output format options:

- c, --csv**
 Output stats in a csv-like format for parsing
- q, --no-header**
 Don't display column headers/footers

Miscellaneous options:

- log-level=level, -l level**
 Force the log verbosity level (overrides configuration value).
 Allowed values: CRIT, MAJOR, EVENT, VERB, DEBUG, FULL.
- help, -h**
 Display a short help about command line options.
- version, -V**
 Display version info.

4.3 Reports

This command can provide the following reports:

Filesystem content report (--fs-info option)

This displays the number of entries of each type, and their volume stats.

Example of output:

type	,	count,	volume,	avg_size
dir	,	1780074,	8.02 GB,	4.72 KB
file	,	21366275,	91.15 TB,	4.47 MB
symlink	,	496142,	24.92 MB,	53

Total: 23475376 entries, 100399805708329 bytes (91.31 TB)

Entry info (--entry-info option)

Displays all informations stored in database about an entry.

Example of output:

```
id          : [0x200000400:0x16a94:0x0]
parent_id   : [0x200000007:0x1:0x0]
name        : file.1
path updt   : 2013/10/30 10:25:30
path        : /mnt/lustre/file.1
depth       : 0
user        : root
group       : root
size        : 1.42 MB
spc_used    : 1.42 MB
creation    : 2013/10/30 10:07:17
last_access : 2013/10/30 10:15:28
last_mod    : 2013/10/30 10:10:52
type        : file
mode        : rw-r--r--
nlink       : 1
md updt     : 2013/10/30 10:25:30
stripe_cnt, stripe_size, pool: 2, 1.00 MB,
stripes     : ost#1: 30515, ost#0: 30520
```

Fileclasses summary (--class-info option)

This generates a summary of fileclasses used for purge policy.

Example of output:

purge class	,	count,	spc_used,	volume,	min_size,	max_size,	avg_size
BigFiles	,	1103,	19.66 TB,	20.76 TB,	8.00 GB,	512.00 GB,	19.28 GB
EmptyFiles	,	1048697,	7.92 GB,	4.15 GB,	0,	1.96 GB,	4.15 KB
SmallFiles	,	20218577,	9.63 TB,	9.67 TB,	0,	95.71 MB,	513.79 KB
ImportantFiles	,	426427,	60.75 TB,	60.86 TB,	16.00 MB,	7.84 GB,	149.66 MB

Note: if you need to list all entries of a given fileclass, use the '--dump-all' and the '--filter-class' options together:

```
rbh-report --dump-all --filter-class project_B
```

User info report (--user-info option)

This displays about the same statistics as 'fs-info' for each user (or only for the user given in parameter).

Example of output:

user	,	type,	count,	spc_used,	avg_size
foo	,	dir,	75450,	306.10 MB,	4.15 KB
foo	,	file,	116396,	11.14 TB,	100.34 MB

Total: 191846 entries, 12248033808384 bytes used (11.14 TB)

Group info report (--group-info option)

Same report as 'user-info', for groups.

Info per user and group (--split-user-groups option, or briefly: -S)

This option splits user stats by group.

E.g.: `rbh-report -u john -S`

user	,	group,	type,	count,	spc_used,	avg_size
john	,	grp1,	dir,	208,	106.49 KB,	0
john	,	grp1,	file,	259781,	781.21 GB,	3.08 MB
john	,	grp2,	dir,	125,	64.00 KB,	0
john	,	grp2,	file,	34525,	4.26 GB,	123.46 KB

Top directories (--top-dirs option)

This option displays directories with the highest number of child entries.

Useful information is given for each of them: path, owner, number of entries, avg size of dir entries, last modification time.

Example of output:

rank,	path,	dircount,	avgsz,	user,	group,	last_mod
1,	/lustre/90k,	89957,	0,	john,	pr01,	2012/03/21 21:50:14
2,	/lustre/foooo,	80939,	79.13 KB,	tod,	pr32,	2011/05/25 13:33:05
3,	/lustre/results,	73983,	197.49 KB,	eddy,	pr32,	2011/06/15 19:23:40
4,	/lustre/bar,	73524,	1.72 MB,	charlie,	pr87,	2012/09/15 19:36:17
5,	/lustre/log,	53693,	350.94 KB,	lily,	pr14,	2012/10/22 18:32:31

- By default, directories are sorted by entry count
- '--by-avgsz' option sorts dirs by average file size
- '--reverse' option reverses sort order (e.g. smallest first)
- Use '--count-min N' option to only display directories with at least N entries.

Example of use: find directories with plenty of small files

(directories with at least 1000 entries, sorted by average file size, from the smaller avgsz to the bigger)

`rbh-report --topdirs --count-min=1000 --by-avgsz --reverse`

rank,	path,	dircount,	avgsz,	user,	group,	last_mod
1,	/lustre/dir1,	1121,	6.01 KB,	foo,	gp1,	2012/04/29 18:20:34
2,	/lustre/dir2,	1543,	7.75 KB,	foo,	gp1,	2012/04/17 15:34:55
3,	/lustre/x,	1029,	12.88 KB,	bar,	gp2,	2011/02/16 01:22:57
4,	/lustre/y,	1019,	12.99 KB,	bar,	gp2,	2011/01/21 11:16:06
...						

Top file size (--top-size option)

This option displays a list of largest files, with useful information: path, size, last access time, last modification time, owner, stripe information.

Example of output:

rank,	path,	size,	user,	group,	last_access,	last_mod,	purge class
1,	/tmp/file.big1,	512.00 GB,	foo1,	p01,	2012/10/14 17:41:38,	2011/05/25 14:22:41,	BigFiles
2,	/tmp/file2.tar,	380.53 GB,	foo2,	p01,	2012/10/14 21:38:07,	2012/02/01 14:30:48,	BigFiles
3,	/tmp/big.1,	379.92 GB,	foo1,	p02,	2012/10/14 20:24:20,	2012/05/17 17:40:57,	BigFiles
...							

Top disk space consumers (--top-users option)

Display users who consume the larger disk space.

rank,	user	, spc_used,	count,	avg_size
1,	usr0021	, 11.14 TB,	116396,	100.34 MB
2,	usr3562	, 5.54 TB,	575,	9.86 GB
3,	usr2189	, 5.52 TB,	9888,	585.50 MB
...				

Useful options:

- ‘--by-count’ option sorts users by entry count
- ‘--by-avgsize’ option sorts users by average file size
- ‘--reverse’ option reverses sort order (e.g. smallest first)
- Use ‘--count-min *N*’ option to only display users with at least *N* entries.
- ‘--by-size-ratio’ option makes it possible to sort users using the percentage of files in the given range.

Example of use: find users with the highest percentage of files < 32MB

(users sorted by file percentage in the range [0-32M[, with at least 1000 entries)

`rbh-report --topusers --by-szratio=0..31M --count-min=1000`

rank,	user	, spc_used,	count,	avg_size,	...	ratio(0..31M)
1,	foo1	, 1.23 GB,	2915,	440.57 KB,	...	100.00%
2,	bar2	, 1.60 MB,	3876,	42.50 KB,	...	100.00%
3,	foo2	, 511.36 MB,	11298,	47.79 KB,	...	99.95%
4,	bar3	, 40.06 GB,	5247,	39.35 MB,	...	99.77%
5,	usr202,	8.72 GB,	58152,	5.10 MB,	...	99.75%

Size profile (--size-profile option)

This option can be used to display the file size profile in the filesystem (with `--fs-info`), for a user (with `--user-info`), or a group (with `--group-info`).

`rbh-report -u foo --szprof`

user,	type,	count,	spc_used,	avg_size,	0,	1~31,	32~1K-,	1K~31K,	32K~1M-,	1M~31M,	32M~1G-,	1G~31G,	32G~1T-,	+1T
foo	, file,	116396,	11.14 TB,	100.34 MB,	56,	2536,	28661,	11054,	16286,	14443,	43291,	68,	1,	0

Note: such information is also available as chart via the robinhood’s web UI

robinhood activity (--activity option)

This reports the last actions Robinhood did, and their status: last filesystem scan, last purge...

Filesystem scan activity:

Current scan interval: 2.8d

Previous filesystem scan:

start: 2012/02/07 11:16:12

```

        duration:          2h 28min 50s

Last filesystem scan:
    status:          running
    start:           2012/02/08 15:00:26 (1h 13min 16s ago)
    last action:     2012/02/08 16:00:27 (1min 15s ago)

Statistics:
    entries scanned: 3757494
    errors:          0
    timeouts:        0
    # threads:       1
    average speed:   4177 entries/sec
>>> current speed: 1985 entries/sec

```

Changelog stats:

```

Last read record id:      71812545
Last record read time:    2012/02/08 16:00:14
Last committed record id: 71812545
Changelog stats:
    type          total (diff) (rate)
    MARK:         135406 (+28)  (0.03/sec)
    CREAT:        16911356 (+82)  (0.09/sec)
    MKDIR:         629730
    HLINK:         128048
    SLINK:         601645 (+11)  (0.01/sec)
    MKNOD:          0
    UNLNK:        11150277 (+47)  (0.05/sec)
    RMDIR:         411672
    RNMFM:        1566093
    RNMT0:        1572679 (+26)  (0.03/sec)
    OPEN:          0
    CLOSE:         0
    IOCTL:         0
    TRUNC:         0
    SATTR:        6674152 (+24)  (0.03/sec)
    XATTR:         0
    HSM:           0
    MTIME:        8201315 (+193) (0.21/sec)
    CTIME:         0
    ATIME:         0

```

Storage unit usage max: 55.33%

```

Last purge: 2011/05/04 13:45:44
Target: OST #16
Status: OK (83 files purged, 958MB)

```

Dump commands (--dump, --dump-user, --dump-group, --dump-ost)

These options can be used for listing entries with a given criteria.
They can be used with filtering options on path (-P) or fileclass (-C).

Example: listing all entries of user 'foo' in the 'BigFiles' class:

```

# rbh-report --dump-user foo -C BigFiles

type,      size,      user,      group,      purge class, path
file, 32.00 GB,      foo,  grp2191,      BigFiles, /lustre/foo/job2389/data.1123
file, 135.32 GB,      foo,  grp2191,      BigFiles, /lustre/foo/job1223/data.2332

```

Total: 2 entries, 179658481991 bytes (167.32 GB)

[new 2.5] Example: listing all entries in OSTs 2 + 5 to 8 (using OST sets)

Note that a file may be striped on an OST, but may have no data on it (depending on file size, stripe order and stripe size). This is indicated by the last column in the report:

```

# rbh-report --dump-ost 2,5-8

```

```

type,      size,      path, stripe_cnt, stripe_size, stripes, data_on_ost[2,5-8]
file,  8.00 MB,  /fs/dir.1/file.8, 2, 1.00 MB, ost#2: 797094, ost#0: 796997, yes
file, 29.00 MB, /fs/dir.1/file.29, 2, 1.00 MB, ost#2: 797104, ost#0: 797007, yes
file,  1.00 MB,  /fs/dir.4/file.1, 2, 1.00 MB, ost#3: 797154, ost#2: 797090, no
file, 27.00 MB, /fs/dir.1/file.27, 2, 1.00 MB, ost#3: 797167, ost#2: 797103, yes
file, 14.00 MB, /fs/dir.5/file.14, 2, 1.00 MB, ost#3: 797161, ost#2: 797097, yes
file, 13.00 MB, /fs/dir.7/file.13, 2, 1.00 MB, ost#2: 797096, ost#0: 796999, yes
file, 24.00 KB, /fs/dir.1/file.24, 2, 1.00 MB, ost#1: 797102, ost#2: 797005, no

```

5 rbh-find

5.1 Overview

This is a clone of the standard ‘find’ command, much faster, as is it based on robinhood’s database (a database is much more adapted for performing queries based on criteria, than a filesystem which is more adapted for IOs).

If you are using Lustre v2 Changelogs, you will get an even more fresh result as the database is fed in soft real-time, whereas the long time of running standard ‘find’ results in an out-of-date result at the end.

Note: this command may require an access to the filesystem (not only a connection to the DB).

5.2 Syntax

* = new in version 2.5.1

Usage: rbh-find [options] [path|fid]...

Filters:

```

-user user
-group group
* -nouser
* -nogroup
-type type
  'f' (file), 'd' (dir), 'l' (symlink), 'b' (block), 'c' (char), 'p' (named pipe/FIFO),
  's' (socket)
-size size crit
  [-|+]<val>[K|M|G|T]
-name filename
-mtime time crit
  [-|+]<val>[s|m|h|d|y] (s: sec, m: min, h: hour, d:day, y:year. default unit is days)
-crttime time crit
  [-|+]<val>[s|m|h|d|y] (s: sec, m: min, h: hour, d:day, y:year. default unit is days)
-mmin minute crit
  same as '-mtime Nm'
-msec second crit
  same as '-mtime Ns'
-atime time crit
  [-|+]<val>[s|m|h|d|y] (s: sec, m: min, h: hour, d:day, y:year. default unit is days)
-amin minute crit
  [-|+]<val>[s|m|h|d|y] (s: sec, m: min, h: hour, d:day, y:year. default unit is days)
-ost ost_index
* -pool ost_pool

-not, -!  Negate the next option

```

Output options:

```

-ls          Display attributes
* -lsost     Display OST information
* -print     Display the fullpath of matching entries (this is the default, unless -ls,
            -lsost or -exec are used).

```

Actions:

```

* -exec "cmd"
  Execute the given command for each matching entry. Unlike classical 'find',
  cmd must be a single (quoted) shell param, not necessarily terminated with ';'.
  '{}' is replaced by the entry path. Example: -exec 'md5sum {}'

```

Behavior:*** -nobulk**

When running `rbh-find` on the filesystem root, `rbh-find` automatically switches to bulk DB request instead of browsing the namespace from the DB. This speeds up the query, but this may result in an arbitrary output ordering, and a single path may be displayed in case of multiple hardlinks. Use `-nobulk` to disable this optimization.

Program options:

-f `config_file`
-d `log_level`
 CRIT, MAJOR, EVENT, VERB, DEBUG, FULL
-h, --help
 Display a short help about command line options.
-V, --version
 Display version info

Notice:

- The path or fid argument is optional. If not specified, the command will run on the entire filesystem specified in robinhood's config file.
- Size can be specified with a suffix (eg. 10M for 10 MB). No suffix means 'bytes'.
- Time criteria can be suffixed (e.g. 10h for 10 hours). No suffix means 'days', like the standard *find* syntax).
- Specifying '+' before a value matches values strictly higher than the value (e.g. +10M matches files whose size is > 10MB).
- Specifying '-' before a value matches values strictly smaller than the value (e.g. -10M matches files whose size is < 10MB).
- For performance reasons, *atime* value stored in robinhood DB may be outdated. So the returned values for "-atime" and "-amin" options may not be accurate. Make sure to double check *atime* value in filesystem for exact matching.
- **[new 2.5.1]** When running `rbh-find` on the filesystem root, `rbh-find` automatically switches to bulk DB request instead of browsing the namespace from the DB. This speeds up the query, but this may result in an arbitrary output ordering, and a single path may be displayed in case of multiple hardlinks. You need to specify the **-nobulk** option to disable this optimization.

5.3 Performance

Performance comparison on a 1 million entry Lustre filesystem:

find

```
find /lustre -user foo -type f -size -32M -ls
(no possible criteria on OST index)
> 58m13s
```

lfs find

```
lfs find /lustre -user foo -type f --obd lustre-OST0001
> 20m46s
(much longer if a criteria on size is specified)
```

rbh-find

```
rbh-find -user foo -type f -size -32M -ost 1 -ls
```

(criteria on both size and OST index)
> 1.2s

⇒ ~3000 times faster than find

6 rbh-du

6.1 Overview

This is a clone of the standard 'du' command, querying robinhood's DB, with additional features and enhancements:

- filtering per user (-u option), per group (-g option) or per type (-t option)
- more detailed outputs (-d option): display entry count, size and disk usage per type.

Note: this command requires an access to the filesystem (not only a connection to the DB).

6.2 Syntax

Usage: rbh-du [options] [path|fid]

Filters:

- u user
- g group
- t type
'f' (file), 'd' (dir), 'l' (symlink), 'b' (block), 'c' (char),
'p' (named pipe/FIFO), 's' (socket)

Output options:

- s, --sum
display total instead of stats per argument
- c, --count
display entry count instead of disk usage
- b, --bytes
display size instead of disk usage (display in bytes)
- k, --kilo
display disk usage in blocks of 1K (default)
- m, --mega
display disk usage in blocks of 1M
- H, --human-readable
display in human readable format (e.g 512K 123.7M)
- d, --details
show detailed stats: type, count, size, disk usage
(display in bytes by default)

Program options:

- f config file
- l log level
- h, --help
Display a short help about command line options.
- V, --version
Display version info

6.3 Performance

du

```
du -sh /lustre/grp1234/usr123
12T      /lustre/grp1234/usr123
> 3m22.108s
```

rbh-du

```
rbh-du -Hd /lustre/grp1234/usr123
dir count:75448, size:305.8M, spc_used:306.1M
file count:116250, size:11.1T, spc_used:11.1T
> 16.815s
```

⇒ 12 times faster than *du* (for 200k entries in user's directory)

7 [new 2.5] rbh-diff

7.1 Overview

rbh-diff command performs a filesystem scan and displays differences with the information currently stored in robinhood database. Optionally, it can apply those changes to the database or revert the detected changes in the filesystem.

It can be used:

- for disaster recovery purpose: after restoring an outdated version of filesystem metadata from a snapshot, **rbh-diff** can restore the metadata changes that occurred after the snapshot time.
Moreover, with a Lustre 2.1 filesystem, it is able to restore newly created files and link them to their data on OSTs. It can also rebuild a MDS from scratch in case of disaster! For more details, see the *Lustre disaster recovery guide* in robinhood documentation repository.
- if you are just curious about what changed in the filesystem since the last robinhood scan.
- for debugging purpose: to detect robinhood database inconsistencies.

7.2 Syntax

rbh-diff

List differences between robinhood database and the filesystem.

Options:

- s dir, --scan=dir
Only scan the specified subdir.
- d attrset, --diff=attrset :
Display changes for the given set of attributes.
attrset is a list of options in:
path, posix, stripe, all, notimes, noatime.
- a {fs|db}, --apply[={fs|db}]
db (default): apply changes to the database using the filesystem as the reference.
fs: revert changes in the filesystem using the database as the reference.

--dry-run
If --apply=fs, display operations on filesystem without performing them.

-f file, --config-file=file
Path to configuration file (or short name).

-l level, --log-level=level
Force the log verbosity level (overrides configuration value).
Allowed values: CRIT, MAJOR, EVENT, VERB, DEBUG, FULL.

-h, --help
Display a short help about command line options.

-V, --version
Display version info

7.3 Output overview

The example below simply lists the changes in the filesystem since the last scan.

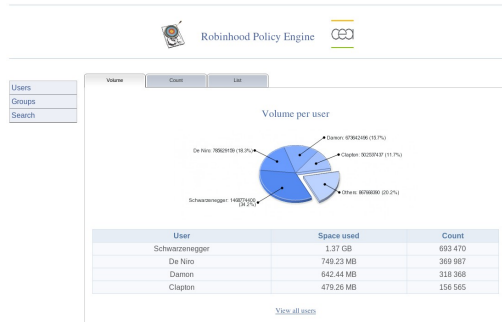
- **Changes** are displayed as 2 lines:
 - the first line, starting with '-', indicates the old attribute value (from robinhood DB)
 - the second line, stating with '+', indicated the new attribute value (from the filesystem)
- **New entries** in the filesystem are represented in lines starting with '++'.
- **Removed entries** are represented in lines starting with '--'.

```
# rbh-diff
---db
+++fs
-[0x200000bd0:0x1a:0x0] path='/mnt/lustre/file'
+[0x200000bd0:0x1a:0x0] path='/mnt/lustre/fname'
-[0x200000bd0:0x16:0x0] size=0,blocks=0
+[0x200000bd0:0x16:0x0] size=18,blocks=1
++[0x200000bd0:0xe:0x0] path='/mnt/lustre/dir.1/file.new',parent=[0x200000bd0:0x1:0x0],
    type=file,nlink=1,mode=0644,owner=root,group=root,size=0,blocks=0,depth=1,
    access=1383128424,modif=1383128424,creation=1383128424,stripes={ost#0:35561,
    ost#1:35497},stripe_count=2,stripe_size=1048576
++[0x200000bd0:0xf:0x0] path='/mnt/lustre/dir.new',parent=[0x61ab:0x2dd8dd16:0x0],
    type=dir,nlink=2,mode=0755,owner=root,group=root,size=4096,blocks=8,depth=0,dircount=0,
    access=1383128424,modif=1383128424,creation=1383128424 -[0x200000bd0:0x12:0x0] owner=root
-[0x200000bd0:0x17:0x0] path='/mnt/lustre/dir.2/d',parent=[0x200000bd0:0x12:0x0]
+[0x200000bd0:0x17:0x0] path='/mnt/lustre/dir.1/d',parent=[0x200000bd0:0x11:0x0]
-[0x200000bd0:0x11:0x0] mode=0750
+[0x200000bd0:0x11:0x0] mode=0700
--[0x200000bd0:0x15:0x0] path=/mnt/lustre/dir.1/b
```

8 Web interface

8.1 Overview

Web interface makes it possible for an administrator to visualize top disk space consumers (per user or per group), top inode consumers with fancy charts, details for each user. It also makes it possible to search for specific entries in the filesystem.



Group	Type	Blocks	Size	Count
Gouverneur	file	2083728	807761888	335466
	dir	91856	47030272	11482
Acteur	file	2678880	686194688	334860
	dir	93296	47767052	11062

8.2 Installation and configuration

You can install it on any machine with a web server (not necessarily the robinhood or the database node). Of course, the web server must be able to contact the Robinhood database.

Requirements: php/mysql support must be enabled in the web server configuration. The following packages must be installed on the web server: php, php-mysql, php-xml, php-pdo, php-gd

The following parameter must be set in httpd.conf:
AllowOverride All

Install robinhood interface:

- install robinhood-webgui RPM on the web server (it will install php files into /var/www/html/robinhood)
- or
- untar the robinhood-webgui tarball in your web server root directory (e.g. /var/www/http)

Configuration:

In a web browser, enter the robinhood URL: <http://yourserver/robinhood>

The first time you connect to this address, fill-in database parameters (host, login, password, ...).

Those parameters are saved in:

/var/www/http/robinhood/app/config/database.xml

Configuration

Enter database configuration

DBMS:

Host (localhost by default):

Database name:

User name:

Password:

That's done. You can enjoy statistics charts.

9 Configuration and admin helper (rbh-config)

rbh-config is a script (installed in /usr/sbin) that helps administrator for configuration and maintenance operations (mostly on robinhood database). It can be used as an interactive command, or in batch mode:

- Interactive mode: just specify an action, it prompts for additional parameters.
E.g: rbh-config test_db

- Batch mode: specify all required parameters on command line.
E.g.: `rbh-config test_db "robinhood_scratch" "passwd"`

Available actions:

- **precheck_db**: check database packages and service.
- **create_db**: create robinhood database.
- **empty_db**: clear robinhood database content.
- **test_db**: test if the database exists and test login on it.
- **repair_db**: check tables and fix them after a mysql server crash.
- **[new 2.5.3] reset_acct**: to rebuild accounting information in case it looks corrupted (this needs to stop/start robinhood daemon: see `rbh-config` inline help).
- **reset_classes**: reset fileclasses after a change in config file.
- **enable_chglogs**: enable ChangeLogs (Lustre 2.x only, must be executed on MDT).
- **backup_db**: backup robinhood database.
- **[new 2.5] optimize_db**: defragments the database to improve its performance and reducing its disk usage.

10 Performance tunings

10.1 Lustre tunings and workarounds

Several bugs or bad behaviours in Lustre can make your node crash or use a lot of memory when Robinhood is scanning or massively purging entries in the FileSystem. Here are some workarounds we had to apply on our system for making it stable:

- If your system “Oops” in `statahead` function, disable this feature:
`echo 0 > /proc/fs/lustre/llite/*/statahead_max`
- CPU overload and client performance drop when free memory is low (bug #17282):
in this case, `lru_size` must be set at `CPU_count * 100`:
`lctl set_param ldlm.namespaces.*.lru_size=800`

10.2 Linux kernel tunings

Robinhood daemon retrieves attributes for large sets of entries in filesystem:

- When scanning, it needs to retrieve attributes of all objects in the filesystem;
- When purging, it checks the attributes of entries before purging them.

This results in loading many inodes and direntries in Linux VFS cache.

If the free memory of the machine is low, or if the machine swaps, this may be due to this cache. To check this, you can read the `/proc/slabinfo` pseudo-file that reports how many objects are allocated by the system, and their size.

For reducing the size of this cache, you can make VFS garbage collection more aggressive by setting the `/proc/sys/vm/vfs_cache_pressure` parameter. By default, its value is 100. If you increase it, garbage collection will be more aggressive and VFS cache will use less memory.

E.g:
echo 1000 > /proc/sys/vm/vfs_cache_pressure

10.3 Optimize scanning vs reporting speed

By default, robinhood is optimized for speeding up common accounting reports (by user, by group, ...), but this can slow database operations during filesystem scans. If you only need specific reports, you can disable some parameters to make scan faster.

For instance, if you only need usage reports by user, you had better disable *group_acct* parameter; this will improve scan performance. In this case, reports on groups will still be available, but their generation will be slower: if you request a *group-info* report and if *group_acct* is *off*, the program will iterate through all the entries (complexity: $O(n)$ with n =number of entries). If *group_acct* is *on*, robinhood will directly access the data in its accounting table, which is quite instantaneous (complexity: $O(1)$).

Performance example: with *group_acct* parameter activated, *group-info* report is generated in 0.01sec for 1M entries. If *group_acct* is disabled, the same report takes about 10sec.

10.4 Database tunings

For managing very large filesystems, some tuning is needed to optimize database performance and fit with available memory. Of course, using large buffers and memory caches will make DB requests faster, but if buffers are oversized, the DB engine and the client may use too much memory, slow-down filesystem performances or make the machine swap.

MySQL server tuning is to be done in `/etc/my.cnf`.

By default (since version 2.4), Robinhood uses innodb MySQL engine (see database parameters in section 3.13 to use MyIsam instead).

It is recommended to tune the following parameters for innodb:

innodb_file_per_table

50% to 90% of the physical memory
innodb_buffer_pool_size=16G

2*nbr_cpu_cores
innodb_thread_concurrency=32

memory cache tuning
innodb_max_dirty_pages_pct=15

robinhood is massively multithreaded: set enough connections
for its threads, and its multiple instances
max_connections=256

If you get DB connection failures, increase this parameter:
connect_timeout=60

This parameter appears to have a significant impact on performances:
see this article to tune it appropriately:

#<http://www.mysqlperformanceblog.com/2008/11/21/how-to-calculate-a-good-innodb-log-file-size>

innodb_log_file_size=500M

To manage transactions efficiently, innodb needs a storage backend with high IOPS performances. You can monitor you disk stress by running “sar -d” on your DB storage device: if ‘%util’ field is close to 100%, your database rate is limited by disk IOPS. In this case you have the choice between these 2 solutions, depending on how critical is your robinhood DB content:

- Safe (needs specific hardware): put your DB on a SSD device, or use a write-back capable storage that is protected against power-failures. In this case, no DB operation can be lost.

- Cheap (and unsafe): add this tuning to /etc/my.cnf:
innodb_flush_log_at_trx_commit=2

This results in flushing transactions to disk only every second, which dramatically reduce the required IO rate. The risk is to loose the last second of recorded information if the DB host crashes. This is affordable if you use to scan your filesystem (the missing information will be added to the database during the next scan). If you read Lustre changelogs, then you will need to scan your filesystem after a DB server failure.

This little script is also very convenient to analyze your database performance and it often suggests relevant tunings:

<http://mysqltuner.pl>

10.5 DB operation batching vs. multi-threading

By default (since version 2.5), robinhood performs database operation batching to limit the need for IOPS for the DB storage.

Depending on your DB storage, you may want to parallelize DB accesses instead of batching them. To tune this, refer to section 3.16: *Entry processor pipeline options*.

11 Getting help

If you still have questions once you read this guide...

You can find several tips and answers for frequently asked questions in the wiki pages on the robinhood project website:

<http://robinhood.sf.net>

You can also take a look at the archive of support mailing list on *sourceforge*:

<http://sourceforge.net/projects/robinhood>

→ Mailing Lists

→ robinhood-support: archive / search

If that didn't help, send your question to the support mailing list:

robinhood-support@lists.sf.net